

[illegible]

# SPRING AI

## GENERATIVE ARTIFICIAL INTELLIGENCE CON JAVA

Simone Scannapieco

Corso avanzato per Venis S.p.A, Venezia, Italia

Novembre 2025

### Note

This image shows a single sheet of white paper with horizontal blue or grey ruling lines. The lines are evenly spaced and run across the width of the page, providing a template for handwriting practice or general writing. There are no margins, text, or other markings on the paper.

This image shows a blank sheet of white paper with horizontal ruling lines. The lines are evenly spaced and extend across the width of the page. There are no margins, text, or other markings on the paper.

- ❓ ... ma come usarli per *task* specifici o con conoscenza che a loro manca?!

[illegible]



## This image shows a blank sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

- ➔ Creare una architettura neurale da zero...
- ➔ ... oppure scegliere una architettura in letteratura (per i meno sadici)
- ➔ Addestramento da zero (a partire da pesi e *bias random*)

- ➔ Sfruttare una rete neurale già addestrata su un altro insieme di dati di addestramento
- ➔ Modificare solo alcuni strati (solitamente gli ultimi) per addestrare la rete per i propri scopi

<i><b>Computer Vision</b></i>	<i><b>Full learning</b></i>	<i><b>Transfer learning</b></i>
<b>Numero dati addestramento</b>	$10^3-10^6$	$10^2$
<b>Computazione</b>	Intensiva (GPU)	Media (CPU-GPU)
<b>Tempo di addestramento</b>	Giorni-settimane	Ore-giorni
<b>Accuratezza del modello</b>	Alta	Variabile

### Note

This image shows a blank sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.


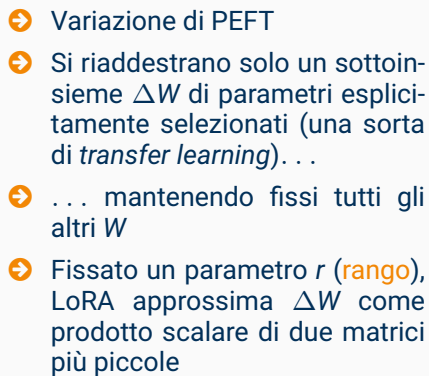
-  Simone Scannapieco

[illegible]

-  Simone Scannapieco

[illegible]





### Note

[illegible]

Il *prodotto scalare* di matrice  $(n \times r)$   $A = \begin{bmatrix} a_{11} & a_{12} & \dots & a_{1r} \\ \vdots & \vdots & \vdots & \vdots \\ a_{n1} & a_{n2} & \dots & a_{nr} \end{bmatrix}$  e matrice  $(r \times m)$

$$B = \begin{bmatrix} b_{11} & b_{12} & \dots & b_{1m} \\ \vdots & \vdots & \vdots & \vdots \\ b_{r1} & b_{r2} & \dots & b_{rm} \end{bmatrix} \quad \text{é la matrice } (n \times m)$$

$$A \cdot B = \begin{bmatrix} a_{11} * b_{11} + \dots + a_{1r} * b_{r1} & \dots & a_{11} * b_{1m} + \dots + a_{1r} * b_{rm} \\ \vdots & \vdots & \vdots \\ a_{n1} * b_{11} + \dots + a_{nr} * b_{r1} & \dots & a_{n1} * b_{1m} + \dots + a_{nr} * b_{rm} \end{bmatrix}$$

- ➔ In pratica, fissato  $r$ , LoRA computa  $A$  e  $B$  tale per cui  $\Delta W = A \cdot B$
- ➔ Ma perché é così potente?!

### Note

This image shows a blank sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

$$\Delta W = \begin{bmatrix} 5 & 1 & -1 & 3 & 4 \\ 15 & 3 & -3 & 9 & 12 \\ 35 & 7 & -7 & 21 & 28 \\ -20 & -4 & 4 & -12 & -16 \\ 10 & 2 & -2 & 6 & 8 \end{bmatrix} \xrightarrow{\text{LoRA}(r=1)} A = \begin{bmatrix} 1 \\ 3 \\ 7 \\ -4 \\ 2 \end{bmatrix}, B = [5 \quad 1 \quad -1 \quad 3 \quad 4]$$

-  Simone Scannapieco

### Note

This image shows a blank sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.


 Simone Scannapieco
  Spring AI - Corso avanzato
  Venis S.p.A, Venezia, IT
 11 / 19

## This image shows a blank sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

- ➔ Quantizzazione a `float16` e `bfloat16` usati maggiormente per addestramento
- ➔ `bfloat16` generalmente preferito a `float16`
- ➔ Addestramento a `float32` riservato alle *big companies*
- ➔ Quantizzazioni inferiori disponibili (`int8`, `int4`), ma consigliate per inferenza

## Note

This image shows a blank sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

This image shows a blank sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

- | Formato | Riduzione memoria | Uso principale  | Accuratezza |
|---------|-------------------|-----------------|-------------|
| int8    | ~50%              | Inferenza       | Alta        |
| int4    | ~75%              | Inferenza       | Media-Alta  |
| GPTQ    | ~75%              | Inferenza (GPU) | Alta        |
| GGUF    | 50-80%            | Inferenza (CPU) | Variabile   |

[illegible]



- 15 / 19

This image shows a single sheet of white paper with horizontal blue or grey ruling lines. The lines are evenly spaced and run across the width of the page, leaving small margins at the top and bottom. There are no vertical margin lines, and the paper is otherwise completely blank.



-  Simone Scannapieco

[illegible]

- The screenshot displays the 'LLM Inference: VRAM & Performance Calculator' interface. The left sidebar contains configuration options for Inference, Fine-tuning, KV Cache Quantization, Inference Quantization, Hardware Configuration, Batch Size, Segment Length, and Concurrent Users. The main panel shows the selected model 'DeepSeek-R1.5B', quantization settings (FP16), hardware (RTX 3060 (12GB)), and a batch size of 1. The right sidebar, titled 'Performance & Memory Results', features a circular progress indicator showing 62.7% usage, a 'MODERATE' performance level, and a memory requirement of 7.52 GB. Below this, it lists generation speed (~38 tokens/sec), time to first token (~223ms), total throughput (~38 tokens/sec), and provides download links for DeepSeek-R1.5B in various quantization formats (FP16, GPTQ, GGUF, EXL2).

Configuration	Value
Select Model	DeepSeek-R1.5B
Inference Quantization	FP16
KV Cache Quantization	FP16/FP2 (Default)
Hardware Configuration	RTX 3060 (12GB)
Batch Size	1
Segment Length	128
Concurrent Users	1

Performance & Memory Results	Value
Usage	62.7%
Performance Level	MODERATE
Memory Requirement	7.52 GB
Generation Speed	~38 tokens/sec
Time to First Token	~223ms
Total Throughput	~38 tokens/sec

17 / 19

## Note

[illegible]



- ➔ Adattamento stile, tono, formato *output*
- ➔ Comportamenti specifici o *task* strutturati
- ➔ *Dataset* 1K–100K esempi, risorse *hardware* limitate
- ➔ Necessità di gestire multipli adattatori per *task* diversi

## This image shows a blank sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.